

Retail Analysis with Walmart Data

March 9, 2024

```
[2]: # Import necessary libraries
import pandas as pd
import seaborn as sns
import numpy as np
import matplotlib.pyplot as plt
from matplotlib import dates
from datetime import datetime
```

0.0.1 Data Understanding

There are sales data available for 45 stores of Walmart in [Kaggle](#). This is the data that covers sales from 2010-02-05 to 2012-11-01.

The data contains these features: - Store - the store number - Date - the week of sales - Weekly_Sales - sales for the given store - Holiday_Flag - whether the week is a special holiday week 1 – Holiday week 0 – Non-holiday week - Temperature - Temperature on the day of sale - Fuel_Price - Cost of fuel in the region - CPI – Prevailing consumer price index - Unemployment - Prevailing unemployment rate

```
[3]: # Load dataset
data = pd.read_csv('Walmart_Store_sales.csv')
data
```

```
[3]:
```

	Store	Date	Weekly_Sales	Holiday_Flag	Temperature	Fuel_Price	\
0	1	05-02-2010	1643690.90	0	42.31	2.572	
1	1	12-02-2010	1641957.44	1	38.51	2.548	
2	1	19-02-2010	1611968.17	0	39.93	2.514	
3	1	26-02-2010	1409727.59	0	46.63	2.561	
4	1	05-03-2010	1554806.68	0	46.50	2.625	
...	
6430	45	28-09-2012	713173.95	0	64.88	3.997	
6431	45	05-10-2012	733455.07	0	64.89	3.985	
6432	45	12-10-2012	734464.36	0	54.47	4.000	
6433	45	19-10-2012	718125.53	0	56.47	3.969	
6434	45	26-10-2012	760281.43	0	58.85	3.882	
		CPI	Unemployment				
0		211.096358	8.106				
1		211.242170	8.106				

2	211.289143	8.106
3	211.319643	8.106
4	211.350143	8.106
...
6430	192.013558	8.684
6431	192.170412	8.667
6432	192.327265	8.667
6433	192.330854	8.667
6434	192.308899	8.667

[6435 rows x 8 columns]

0.0.2 Data Preparation

```
[4]: # Convert date to datetime format and show dataset information
data['Date'] = pd.to_datetime(data['Date'])
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 6435 entries, 0 to 6434
Data columns (total 8 columns):
Store                6435 non-null int64
Date                 6435 non-null datetime64[ns]
Weekly_Sales         6435 non-null float64
Holiday_Flag         6435 non-null int64
Temperature           6435 non-null float64
Fuel_Price           6435 non-null float64
CPI                  6435 non-null float64
Unemployment         6435 non-null float64
dtypes: datetime64[ns](1), float64(5), int64(2)
memory usage: 402.3 KB
```

```
[5]: # checking for missing values
data.isnull().sum()
```

```
[5]: Store                0
Date                    0
Weekly_Sales           0
Holiday_Flag           0
Temperature             0
Fuel_Price             0
CPI                    0
Unemployment           0
dtype: int64
```

```
[6]: # Splitting Date and create new columns (Day, Month, and Year)
data["Day"] = pd.DatetimeIndex(data['Date']).day
data['Month'] = pd.DatetimeIndex(data['Date']).month
```

```
data['Year'] = pd.DatetimeIndex(data['Date']).year
data
```

```
[6]:
```

	Store	Date	Weekly_Sales	Holiday_Flag	Temperature	Fuel_Price	\
0	1	2010-05-02	1643690.90	0	42.31	2.572	
1	1	2010-12-02	1641957.44	1	38.51	2.548	
2	1	2010-02-19	1611968.17	0	39.93	2.514	
3	1	2010-02-26	1409727.59	0	46.63	2.561	
4	1	2010-05-03	1554806.68	0	46.50	2.625	
...	
6430	45	2012-09-28	713173.95	0	64.88	3.997	
6431	45	2012-05-10	733455.07	0	64.89	3.985	
6432	45	2012-12-10	734464.36	0	54.47	4.000	
6433	45	2012-10-19	718125.53	0	56.47	3.969	
6434	45	2012-10-26	760281.43	0	58.85	3.882	

	CPI	Unemployment	Day	Month	Year
0	211.096358	8.106	2	5	2010
1	211.242170	8.106	2	12	2010
2	211.289143	8.106	19	2	2010
3	211.319643	8.106	26	2	2010
4	211.350143	8.106	3	5	2010
...
6430	192.013558	8.684	28	9	2012
6431	192.170412	8.667	10	5	2012
6432	192.327265	8.667	10	12	2012
6433	192.330854	8.667	19	10	2012
6434	192.308899	8.667	26	10	2012

[6435 rows x 11 columns]

0.0.3 Q1: Which store has minimum and maximum sales?

```
[7]: plt.figure(figsize=(15,7))

# Sum Weekly_Sales for each store, then sortded by total sales
total_sales_for_each_store = data.groupby('Store')['Weekly_Sales'].sum().
    ↪sort_values()
total_sales_for_each_store_array = np.array(total_sales_for_each_store) #
    ↪convert to array

# Assigning a specific color for the stores have the lowest and highest sales
clrs = ['lightsteelblue' if ((x < max(total_sales_for_each_store_array)) and (x
    ↪min(total_sales_for_each_store_array))) else 'midnightblue' for x in
    ↪total_sales_for_each_store_array]
```

```

ax = total_sales_for_each_store.plot(kind='bar',color=clrs);


# store have minimum sales
p = ax.patches[0]
print(type(p.get_height()))
ax.annotate("The store has minimum sales is 33 with {0:.2f} $".format((p.
    ↪get_height()))), xy=(p.get_x(), p.get_height()), xycoords='data',
    xytext=(0.17, 0.32), textcoords='axes fraction',
    arrowprops=dict(arrowstyle="->", connectionstyle="arc3"),
    horizontalalignment='center', verticalalignment='center')

# store have maximum sales
p = ax.patches[44]
ax.annotate("The store has maximum sales is 20 with {0:.2f} $".format((p.
    ↪get_height()))), xy=(p.get_x(), p.get_height()), xycoords='data',
    xytext=(0.82, 0.98), textcoords='axes fraction',
    arrowprops=dict(arrowstyle="->", connectionstyle="arc3"),
    horizontalalignment='center', verticalalignment='center')

# plot properties
plt.xticks(rotation=0)
plt.ticklabel_format(useOffset=False, style='plain', axis='y')
plt.title('Total sales for each store')
plt.xlabel('Store')
plt.ylabel('Total Sales');

```

<class 'numpy.float64'>

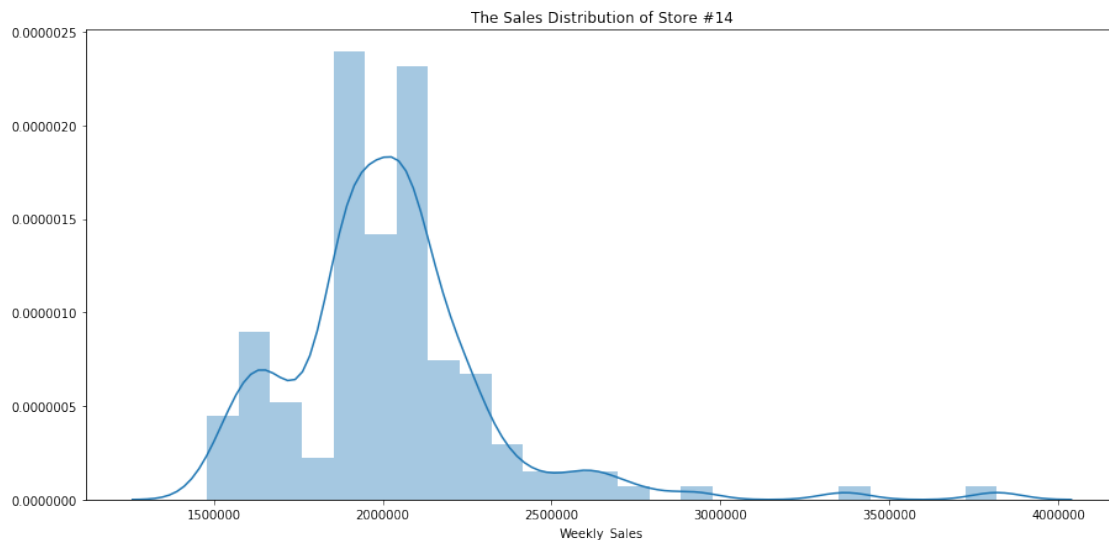


0.0.4 Q2: Which store has maximum standard deviation i.e., the sales vary a lot. Also, find out the coefficient of mean to standard deviation?

```
[8]: # Which store has maximum standard deviation
data_std = pd.DataFrame(data.groupby('Store')['Weekly_Sales'].std().
    ↪sort_values(ascending=False))
print("The store has maximum standard deviation is "+str(data_std.head(1).
    ↪index[0])+" with {0:.0f} $".format(data_std.head(1).Weekly_Sales[data_std.
    ↪head(1).index[0]]))
```

The store has maximum standard deviation is 14 with 317570 \$

```
[9]: # Distribution of store has maximum standard deviation
plt.figure(figsize=(15,7))
sns.distplot(data[data['Store'] == data_std.head(1).index[0]]['Weekly_Sales'])
plt.title('The Sales Distribution of Store #' + str(data_std.head(1).index[0]));
```



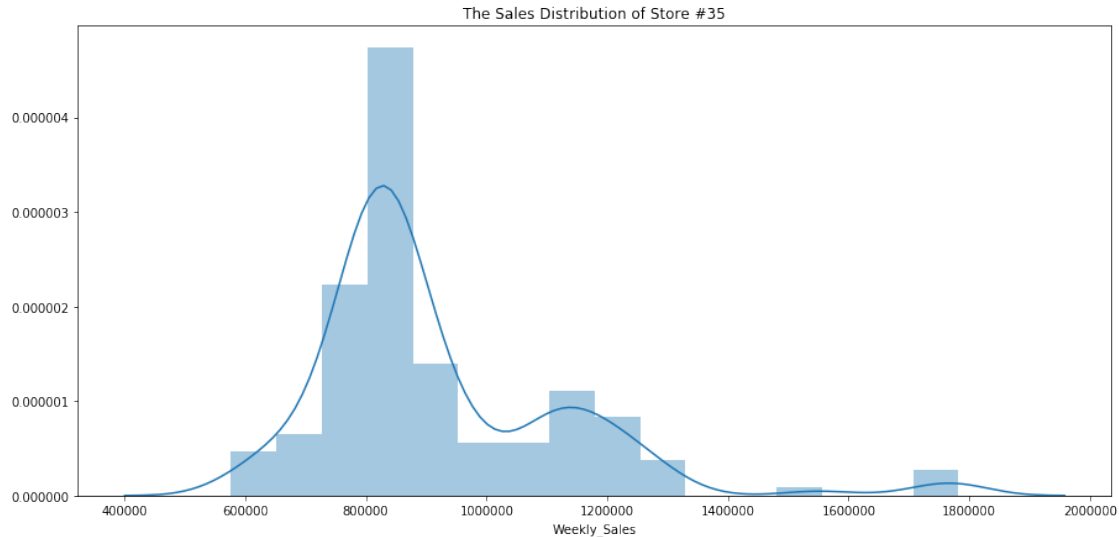
```
[10]: # Coefficient of mean to standard deviation
coef_mean_std = pd.DataFrame(data.groupby('Store')['Weekly_Sales'].std() / data.
    ↪groupby('Store')['Weekly_Sales'].mean())
coef_mean_std = coef_mean_std.rename(columns={'Weekly_Sales': 'Coefficient of_
    ↪mean to standard deviation'})
coef_mean_std
```

```
[10]:      Coefficient of mean to standard deviation
Store
1      0.100292
2      0.123424
3      0.115021
```

4	0.127083
5	0.118668
6	0.135823
7	0.197305
8	0.116953
9	0.126895
10	0.159133
11	0.122262
12	0.137925
13	0.132514
14	0.157137
15	0.193384
16	0.165181
17	0.125521
18	0.162845
19	0.132680
20	0.130903
21	0.170292
22	0.156783
23	0.179721
24	0.123637
25	0.159860
26	0.110111
27	0.135155
28	0.137330
29	0.183742
30	0.052008
31	0.090161
32	0.118310
33	0.092868
34	0.108225
35	0.229681
36	0.162579
37	0.042084
38	0.110875
39	0.149908
40	0.123430
41	0.148177
42	0.090335
43	0.064104
44	0.081793
45	0.165613

```
[11]: # Distribution of store has maximum coefficient of mean to standard deviation
coef_mean_std_max = coef_mean_std.sort_values(by='Coefficient of mean to_
↪standard deviation')
plt.figure(figsize=(15,7))
```

```
sns.distplot(data[data['Store'] == coef_mean_std_max.tail(1).
↪index[0]]['Weekly_Sales'])
plt.title('The Sales Distribution of Store #'+str(coef_mean_std_max.tail(1).
↪index[0]));
```



0.0.5 Q3: Which store/s has good quarterly growth rate in Q3'2012


```
[12]: plt.figure(figsize=(15,7))

# Sales for third quarterly in 2012
Q3 = data[(data['Date'] > '2012-07-01') & (data['Date'] < '2012-09-30')].
↪groupby('Store')['Weekly_Sales'].sum()

# Sales for second quarterly in 2012
Q2 = data[(data['Date'] > '2012-04-01') & (data['Date'] < '2012-06-30')].
↪groupby('Store')['Weekly_Sales'].sum()

# Plotting the difference between sales for second and third quarterly
Q2.plot(ax=Q3.plot('bar',legend=True),kind='bar',color='r',alpha=0.
↪2,legend=True);
plt.legend(["Q3' 2012", "Q2' 2012"]);
```

C:\Users\Zaraah\Anaconda3\lib\site-packages\ipykernel_launcher.py:10:
FutureWarning: `Series.plot()` should not be called with positional arguments,
only keyword arguments. The order of positional arguments will change in the
future. Use `Series.plot(kind='bar')` instead of `Series.plot('bar',)`.
Remove the CWD from sys.path while we load stuff.



```
[13]: # store/s has good quarterly growth rate in Q3'2012 - .
      ↪sort_values(by='Weekly_Sales')
      print('Store have good quarterly growth rate in Q3'2012 is Store '+str(Q3.
      ↪idxmax())+' With '+str(Q3.max())+' $')
```

Store have good quarterly growth rate in Q3'2012 is Store 4 With 25652119.35 \$

0.0.6 Q4: Some holidays have a negative impact on sales. Find out holidays which have higher sales than the mean sales in non-holiday season for all stores together

Holiday Events:

- Super Bowl: 12-Feb-10, 11-Feb-11, 10-Feb-12, 8-Feb-13
- Labour Day: 10-Sep-10, 9-Sep-11, 7-Sep-12, 6-Sep-13
- Thanksgiving: 26-Nov-10, 25-Nov-11, 23-Nov-12, 29-Nov-13
- Christmas: 31-Dec-10, 30-Dec-11, 28-Dec-12, 27-Dec-13

```
[18]: def plot_line(df, holiday_dates, holiday_label):
      fig, ax = plt.subplots(figsize = (15,5))
      ax.plot(df['Date'], df['Weekly_Sales'], label=holiday_label)

      for day in holiday_dates:
          day = datetime.strptime(day, '%d-%m-%Y')
          plt.axvline(x=day, linestyle='--', c='r')

      plt.title(holiday_label)
```



```

x_dates = df['Date'].dt.strftime('%Y-%m-%d').sort_values().unique()
xfmt = dates.DateFormatter('%d-%m-%y')
ax.xaxis.set_major_formatter(xfmt)
ax.xaxis.set_major_locator(dates.DayLocator(1))
plt.gcf().autofmt_xdate(rotation=90)
plt.show()

```

```



total_sales = data.groupby('Date')['Weekly_Sales'].sum().reset_index()
Super_Bowl = ['12-2-2010', '11-2-2011', '10-2-2012']
Labour_Day = ['10-9-2010', '9-9-2011', '7-9-2012']
Thanksgiving = ['26-11-2010', '25-11-2011', '23-11-2012']
Christmas = ['31-12-2010', '30-12-2011', '28-12-2012']

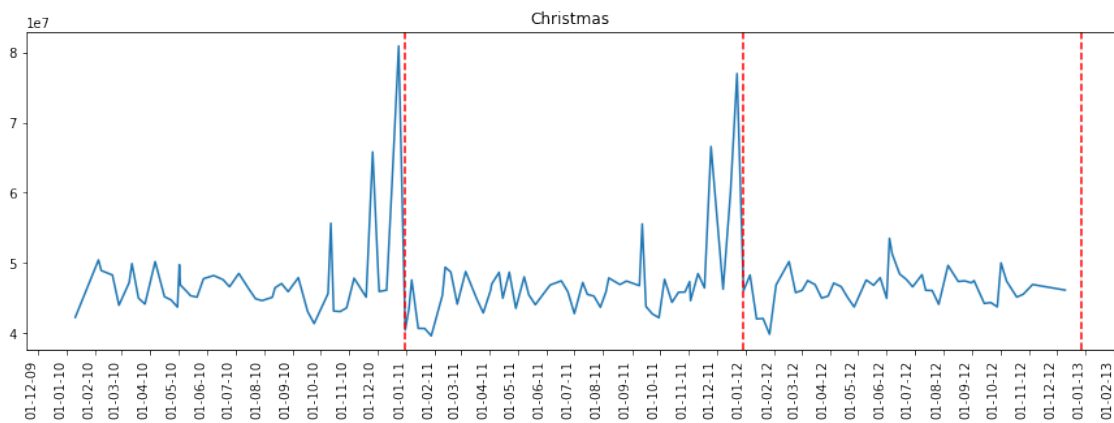

```

```

plot_line(total_sales, Super_Bowl, 'Super Bowl')
plot_line(total_sales, Labour_Day, 'Labour Day')
plot_line(total_sales, Thanksgiving, 'Thanksgiving')
plot_line(total_sales, Christmas, 'Christmas')

```





The sales increased during thanksgiving. And the sales decreased during christmas.

```
[19]: data.loc[data.Date.isin(Super_Bowl)]
```

```
[19]:
```

	Store	Date	Weekly_Sales	Holiday_Flag	Temperature	Fuel_Price	\
1	1	2010-12-02	1641957.44	1	38.51	2.548	
53	1	2011-11-02	1649614.93	1	36.39	3.022	
105	1	2012-10-02	1802477.43	1	48.02	3.409	
144	2	2010-12-02	2137809.50	1	38.49	2.548	
196	2	2011-11-02	2168041.61	1	33.19	3.022	
...	
6202	44	2011-11-02	307486.73	1	30.83	3.034	
6254	44	2012-10-02	325377.97	1	33.73	3.116	
6293	45	2010-12-02	656988.64	1	27.73	2.773	
6345	45	2011-11-02	766456.00	1	30.30	3.239	
6397	45	2012-10-02	803657.12	1	37.00	3.640	



	CPI	Unemployment	Day	Month	Year
1	211.242170	8.106	2	12	2010
53	212.936705	7.742	2	11	2011
105	220.265178	7.348	2	10	2012
144	210.897994	8.324	2	12	2010
196	212.592862	8.028	2	11	2011
...
6202	127.859129	7.224	2	11	2011
6254	130.384903	5.774	2	10	2012
6293	181.982317	8.992	2	12	2010
6345	183.701613	8.549	2	11	2011
6397	189.707605	8.424	2	10	2012


[135 rows x 11 columns]


```
[90]: # Yearly Sales in holidays
Super_Bowl_df = pd.DataFrame(data.loc[data.Date.isin(Super_Bowl)]).
    ↳groupby('Year')['Weekly_Sales'].sum()
Thanksgiving_df = pd.DataFrame(data.loc[data.Date.isin(Thanksgiving)]).
    ↳groupby('Year')['Weekly_Sales'].sum()
Labour_Day_df = pd.DataFrame(data.loc[data.Date.isin(Labour_Day)]).
    ↳groupby('Year')['Weekly_Sales'].sum()
Christmas_df = pd.DataFrame(data.loc[data.Date.isin(Christmas)]).
    ↳groupby('Year')['Weekly_Sales'].sum()

Super_Bowl_df.plot(kind='bar',legend=False,title='Yearly Sales in Super Bowl_
    ↳holiday')
Thanksgiving_df.plot(kind='bar',legend=False,title='Yearly Sales in_
    ↳Thanksgiving holiday')
Labour_Day_df.plot(kind='bar',legend=False,title='Yearly Sales in Labour_Day_
    ↳holiday')
Christmas_df.plot(kind='bar',legend=False,title='Yearly Sales in Christmas_
    ↳holiday')
```

```
[90]: <matplotlib.axes._subplots.AxesSubplot at 0x1b42cb1a08>
```









0.0.7 Q5: Provide a monthly and semester view of sales in units and give insights


```
[101]: # Monthly view of sales for each years
plt.scatter(data[data.Year==2010]["Month"],data[data.
    ↳Year==2010]["Weekly_Sales"])
plt.xlabel("months")
plt.ylabel("Weekly Sales")
plt.title("Monthly view of sales in 2010")
plt.show()

plt.scatter(data[data.Year==2011]["Month"],data[data.
    ↳Year==2011]["Weekly_Sales"])
plt.xlabel("months")
plt.ylabel("Weekly Sales")
plt.title("Monthly view of sales in 2011")
plt.show()

plt.scatter(data[data.Year==2012]["Month"],data[data.
    ↳Year==2012]["Weekly_Sales"])
plt.xlabel("months")
plt.ylabel("Weekly Sales")
plt.title("Monthly view of sales in 2012")
```


```
plt.show()
```






```
[102]: # Monthly view of sales for all years
plt.figure(figsize=(10,6))
plt.bar(data["Month"],data["Weekly_Sales"])
plt.xlabel("months")
plt.ylabel("Weekly Sales")
plt.title("Monthly view of sales")
```

```
[102]: Text(0.5, 1.0, 'Monthly view of sales')
```





```
[110]: # Yearly view of sales
plt.figure(figsize=(10,6))
data.groupby("Year")[["Weekly_Sales"]].sum().plot(kind='bar',legend=False)
plt.xlabel("years")
plt.ylabel("Weekly Sales")
plt.title("Yearly view of sales");
```

<Figure size 720x432 with 0 Axes>



0.1 Build prediction models to forecast demand (Modeling)

```
[111]: # find outliers
fig, axs = plt.subplots(4,figsize=(6,18))
X = data[['Temperature','Fuel_Price','CPI','Unemployment']]
for i,column in enumerate(X):
    sns.boxplot(data[column], ax=axs[i])
```



```
[112]: # drop the outliers
data_new = data[(data['Unemployment']<10) & (data['Unemployment']>4.5) &
↳(data['Temperature']>10)]
data_new
```





```
[112]:
```

	Store	Date	Weekly_Sales	Holiday_Flag	Temperature	Fuel_Price	\
0	1	2010-05-02	1643690.90	0	42.31	2.572	
1	1	2010-12-02	1641957.44	1	38.51	2.548	
2	1	2010-02-19	1611968.17	0	39.93	2.514	
3	1	2010-02-26	1409727.59	0	46.63	2.561	
4	1	2010-05-03	1554806.68	0	46.50	2.625	
...	
6430	45	2012-09-28	713173.95	0	64.88	3.997	
6431	45	2012-05-10	733455.07	0	64.89	3.985	
6432	45	2012-12-10	734464.36	0	54.47	4.000	
6433	45	2012-10-19	718125.53	0	56.47	3.969	
6434	45	2012-10-26	760281.43	0	58.85	3.882	

	CPI	Unemployment	Day	Month	Year
0	211.096358	8.106	2	5	2010
1	211.242170	8.106	2	12	2010
2	211.289143	8.106	19	2	2010
3	211.319643	8.106	26	2	2010
4	211.350143	8.106	3	5	2010
...
6430	192.013558	8.684	28	9	2012
6431	192.170412	8.667	10	5	2012
6432	192.327265	8.667	10	12	2012
6433	192.330854	8.667	19	10	2012
6434	192.308899	8.667	26	10	2012

[5658 rows x 11 columns]

```
[113]: # check outliers
fig, axs = plt.subplots(4,figsize=(6,18))
X = data_new[['Temperature','Fuel_Price','CPI','Unemployment']]
for i,column in enumerate(X):
    sns.boxplot(data_new[column], ax=axs[i])
```



0.1.1 Build Model

```
[121]: # Import sklearn
from sklearn.ensemble import RandomForestRegressor
from sklearn.model_selection import train_test_split
from sklearn import metrics
from sklearn.linear_model import LinearRegression

[122]: # Select features and target
X = data_new[['Store', 'Fuel_Price', 'CPI', 'Unemployment', 'Day', 'Month', 'Year']]
y = data_new['Weekly_Sales']

# Split data to train and test (0.80:0.20)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)


[128]: # Linear Regression model
print('Linear Regression:')
print()
reg = LinearRegression()
reg.fit(X_train, y_train)
y_pred = reg.predict(X_test)
print('Accuracy:', reg.score(X_train, y_train)*100)

print('Mean Absolute Error:', metrics.mean_absolute_error(y_test, y_pred))
print('Mean Squared Error:', metrics.mean_squared_error(y_test, y_pred))
print('Root Mean Squared Error:', np.sqrt(metrics.mean_squared_error(y_test, y_pred)))

sns.scatterplot(y_pred, y_test);
```

Linear Regression:

Accuracy: 12.998176857826893
Mean Absolute Error: 469135.83811432385
Mean Squared Error: 323485572250.08075
Root Mean Squared Error: 568757.9206042591




```
[129]: # Random Forest Regressor
print('Random Forest Regressor:')
print()
rfr = RandomForestRegressor(n_estimators = 400,max_depth=15,n_jobs=5)
rfr.fit(X_train,y_train)
y_pred=rfr.predict(X_test)
print('Accuracy:',rfr.score(X_test, y_test)*100)

print('Mean Absolute Error:', metrics.mean_absolute_error(y_test, y_pred))
print('Mean Squared Error:', metrics.mean_squared_error(y_test, y_pred))
print('Root Mean Squared Error:', np.sqrt(metrics.mean_squared_error(y_test,
    ↪y_pred)))

sns.scatterplot(y_pred, y_test);
```

Random Forest Regressor:

Accuracy: 94.36923602346438
Mean Absolute Error: 72405.98337255421
Mean Squared Error: 20817040949.867416
Root Mean Squared Error: 144281.11778700433



[]:

[]: